

- [Orchestrated Bedrock Pipeline](#)
 - [Overview](#)
 - [Architecture](#)
 - [Orchestrated Workflow](#)
 - [Monitoring & Logging](#)
 - [Retry Mechanism](#)
 - [Usage](#)
 - [What We Are Testing](#)
 - [Test Files](#)
 - [How to Run Tests](#)

Orchestrated Bedrock Pipeline

Overview

This module demonstrates an orchestrated workflow for interacting with Amazon Bedrock (Claude 3.5 Haiku) using a robust, production-ready Python pipeline. It features:

- Modular orchestration
- Monitoring and logging
- Retry mechanisms for reliability
- Testable, extensible design

Architecture

- **orchestrator.py**: Main workflow entrypoint. Handles retries, logging, and error propagation.
- **bedrock_client.py**: Handles all communication with Amazon Bedrock (Claude 3.5 Haiku).
- **logger.py**: Provides rich, timestamped logging for monitoring and debugging.
- **tests/**: Unit and integration tests for pipeline reliability.

Orchestrated Workflow

1. **Input:** User provides a prompt (e.g., "Tell me a joke").
2. **Logging:** All steps are logged with timestamps and color-coded status (info, success, error).
3. **Retry Mechanism:** If Claude invocation fails (network, API, or model error), the pipeline retries up to 3 times with a delay between attempts.
4. **Claude Invocation:** The prompt is sent to Bedrock as a chat message. The response is logged and returned.
5. **Error Handling:** If all retries fail, the error is logged and raised for upstream handling.

Monitoring & Logging

- Uses the **rich** library for colored, timestamped logs.
- Logs every prompt, response, error, and retry attempt.
- Example log output:

```
[INFO 2025-06-21 16:02:20.584230] Received prompt: hi
[ERROR 2025-06-21 16:02:23.474639] Claude invocation failed: ...
[INFO 2025-06-21 16:02:25.633015] Retrying in 2 seconds...
[SUCCESS 2025-06-21 16:02:27.821479] Claude responded successfully
[INFO 2025-06-21 16:02:27.821479] Response: ...
```

Retry Mechanism

- Configurable via the **@retry_on_failure** decorator in **orchestrator.py**.
- Default: 3 attempts, 2 seconds between retries.
- All failures and retries are logged.

Usage

1. Activate your virtual environment and install requirements:

```
pip install -r requirements.txt
```

2. Set up your `.env` file with AWS credentials and region.

3. Run the pipeline:

```
python run_pipeline.py
```

You will be prompted for input, and the orchestrated workflow will handle the rest.

What We Are Testing

- **Pipeline Orchestration:** Ensures the workflow from user prompt to Claude response works as expected.
- **Retry Logic:** Verifies that failed Claude invocations are retried and errors are handled gracefully.
- **Logging:** Confirms that all steps, errors, and retries are logged with correct formatting.
- **Mocked Tests:** Use mocks to simulate Claude responses for fast, reliable unit testing.
- **Integration Test:** Optionally runs a real end-to-end test against Claude 3.5 Haiku (requires valid AWS credentials and Bedrock access).

Test Files

- `tests/test_pipeline.py`: Contains all unit and integration tests for the pipeline.

How to Run Tests

1. **Run all tests:**

```
pytest
```

2. **Check code coverage:**

```
coverage run -m pytest
coverage report -m
```

3. Test details:

- **Mocked tests** (safe for CI): Test pipeline success, empty responses, and error handling using mocked Claude responses.
- **Integration test**: Sends a real prompt to Claude via Bedrock and checks the response (ensure AWS credentials are set).